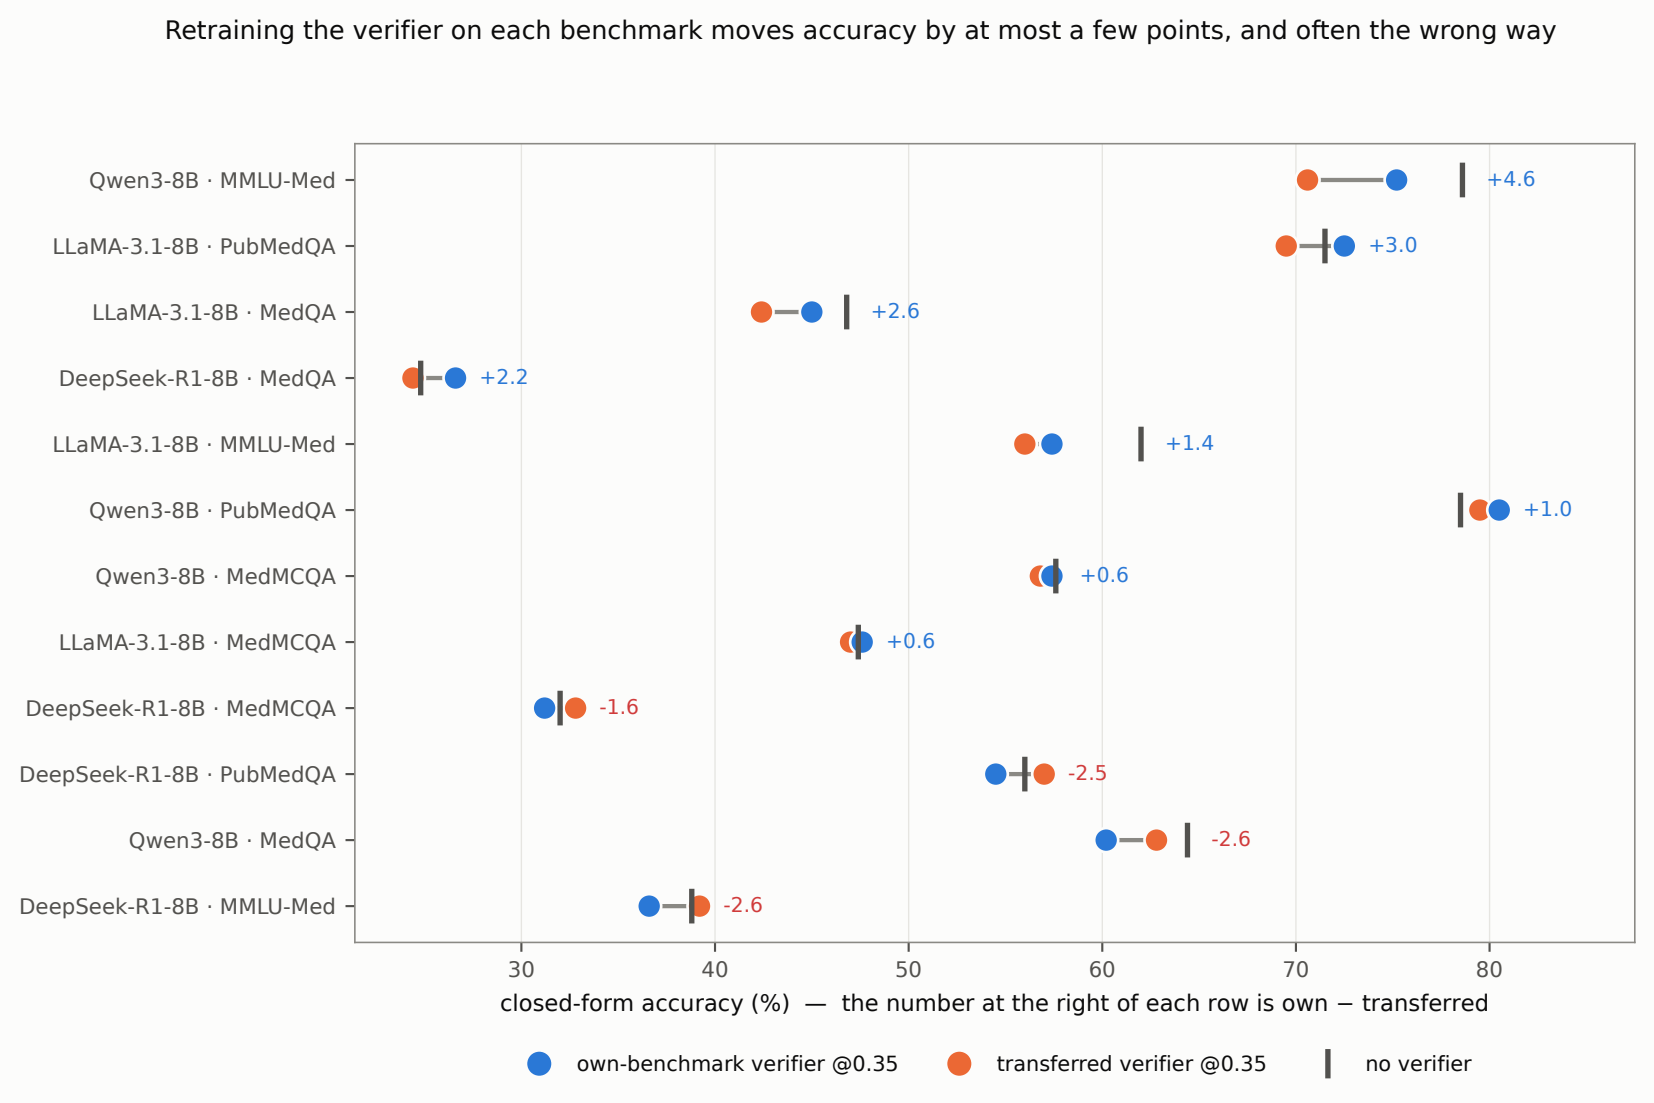
Retraining the verifier on each benchmark moves accuracy by at most a few points, and often the wrong way
Qwen3-8B · MMLU-Med
LLaMA-3.1-8B · PubMedQA
LLaMA-3.1-8B · MedQA
DeepSeek-R1-8B · MedQA
LLaMA-3.1-8B · MMLU-Med
Qwen3-8B · PubMedQA
Qwen3-8B · MedMCQA
LLaMA-3.1-8B · MedMCQA
DeepSeek-R1-8B · MedMCQA
DeepSeek-R1-8B · PubMedQA
Qwen3-8B · MedQA
DeepSeek-R1-8B · MMLU-Med
+4.6
+3.0
+2.6
+2.2
+1.4
+1.0
+0.6
+0.6
-1.6
-2.5
-2.6
-2.6
30
40
50
60
70
80
closed-form accuracy (%) — the number at the right of each row is own − transferred
own-benchmark verifier @0.35
transferred verifier @0.35
no verifier